

02 - Naive Bayesian - Multinomial - Lab

August 26, 2021

0.0.1 === Task ===

- 1) Learn about TfidfVectorizer and replace CountVectorizer with TfidfVectorizer (Explanation Provided in the Lecture)
- 2) Put Multinomial Naive Classification into a class that can transform the data, fit the model and do prediction.
 - In the class, allow users to choose whether to use CountVectorizer or TfidfVectorizer to transform the data.

```
[1]: import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: from sklearn.datasets import fetch_20newsgroups

data = fetch_20newsgroups()
data.target_names
```

```
[2]: ['alt.atheism',
      'comp.graphics',
      'comp.os.ms-windows.misc',
      'comp.sys.ibm.pc.hardware',
      'comp.sys.mac.hardware',
      'comp.windows.x',
      'misc.forsale',
      'rec.autos',
      'rec.motorcycles',
      'rec.sport.baseball',
      'rec.sport.hockey',
      'sci.crypt',
      'sci.electronics',
      'sci.med',
      'sci.space',
      'soc.religion.christian',
      'talk.politics.guns',
      'talk.politics.mideast',
      'talk.politics.misc',
      'talk.religion.misc']
```

```
[3]: categories = ['talk.religion.misc', 'soc.religion.christian',  
                 'sci.space', 'comp.graphics']  
train = fetch_20newsgroups(subset='train', categories=categories)  
test = fetch_20newsgroups(subset='test', categories=categories)
```

```
[4]: print(train.data[0]) #first 300 words  
print("Target: ", train.target[0]) #start with 1, soc.religion.christian
```

From: jono@mac-ak-24.rtsg.mot.com (Jon Ogden)
Subject: Re: Losing your temper is not a Christian trait
Organization: Motorola LPA Development
Lines: 26

In article <Apr.23.02.55.47.1993.3138@geneva.rutgers.edu>, jcyj@tellabs.com
(jcj) wrote:

> I'd like to remind people of the withering of the fig tree and Jesus
> driving the money changers et. al. out of the temple. I think those
> were two instances of Christ showing anger (as part of His human side).
>

Yes, and what about Paul saying:

26 Be ye angry, and sin not: let not the sun go down upon your wrath:
(Ephesians 4:26).

Obviously then, we can be angry w/o sinning.

Jon

Jon Ogden - jono@mac-ak-24.rtsg.mot.com
Motorola Cellular - Advanced Products Division
Voice: 708-632-2521 Data: 708-632-6086

They drew a circle and shut him out.
Heretic, Rebel, a thing to flout.
But Love and I had the wit to win;
We drew a circle and took him in.

Target: 2

```
[5]: #let's use some library to count for us  
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
```

```
[6]: class MyMultiNaive:  
      def __init__(self, method='count'):
```

```

    if (method == 'tfidf'):
        self.vectorizer = TfidfVectorizer()
    else:
        self.vectorizer = CountVectorizer()

def tranform(self, train, test):
    X_train = self.vectorizer.fit_transform(train.data)
    X_test = self.vectorizer.transform(test.data)
    X_test = X_test.toarray()
    y_train = train.target
    y_test = test.target
    return X_train, X_test, y_train, y_test

def likelihood(self, X_class, laplace=1):
    return ((X_class.sum(axis=0)) + laplace) / (np.sum(X_class.sum(axis=0))
↪ + laplace))

def prior(self, X_class, m):
    return X_class.shape[0] / m

def fit(self, X_train, y_train):
    m, n = X_train.shape
    classes = np.unique(y_train) #list of class
    k = len(classes) #number of class

    priors = np.zeros(k) #prior for each classes
    likelihoods = np.zeros((k, n)) #likelihood for each class of each feature

    for idx, label in enumerate(classes):
        X_train_c = X_train[y_train==label]
        priors[idx] = self.prior(X_train_c, m)
        likelihoods[idx, :] = self.likelihood(X_train_c)
    self.priors = priors
    self.likelihoods = likelihoods

def predict(self, X_test, classes):
    return np.log(self.priors) + X_test @ np.log(self.likelihoods.T)

```

0.1 Count

```

[7]: method = 'count' # tfidf, count
    model = MyMultiNaive(method)
    X_train, X_test, y_train, y_test = model.tranform(train, test)
    model.fit(X_train, y_train)

```

```
[8]: classes = np.unique(y_test)
yhat = model.predict(X_test, classes)
yhat = np.argmax(yhat, axis=1)
```

```
[9]: from sklearn.preprocessing import label_binarize
from sklearn.metrics import average_precision_score, classification_report

n_classes = len(np.unique(y_test))

print("Accuracy: ", np.sum(yhat == y_test)/len(y_test))

print("====Average precision score====")
y_test_binarized = label_binarize(y_test, classes=[0, 1, 2, 3])
yhat_binarized = label_binarize(yhat, classes=[0, 1, 2, 3])

for i in range(n_classes):
    class_score = average_precision_score(y_test_binarized[:, i],
    ↪yhat_binarized[:, i])
    print(f"Class {i} score: ", class_score)

print("====Classification report====")
print("Report: ", classification_report(y_test, yhat))
```

```
Accuracy: 0.9168994413407822
====Average precision score====
Class 0 score: 0.9152047938418233
Class 1 score: 0.9069918620723723
Class 2 score: 0.8429395016564877
Class 3 score: 0.7277310085946386
====Classification report====
Report:                precision    recall  f1-score   support

      0         0.95         0.95         0.95         389
      1         0.94         0.96         0.95         394
      2         0.87         0.95         0.91         398
      3         0.92         0.74         0.82         251

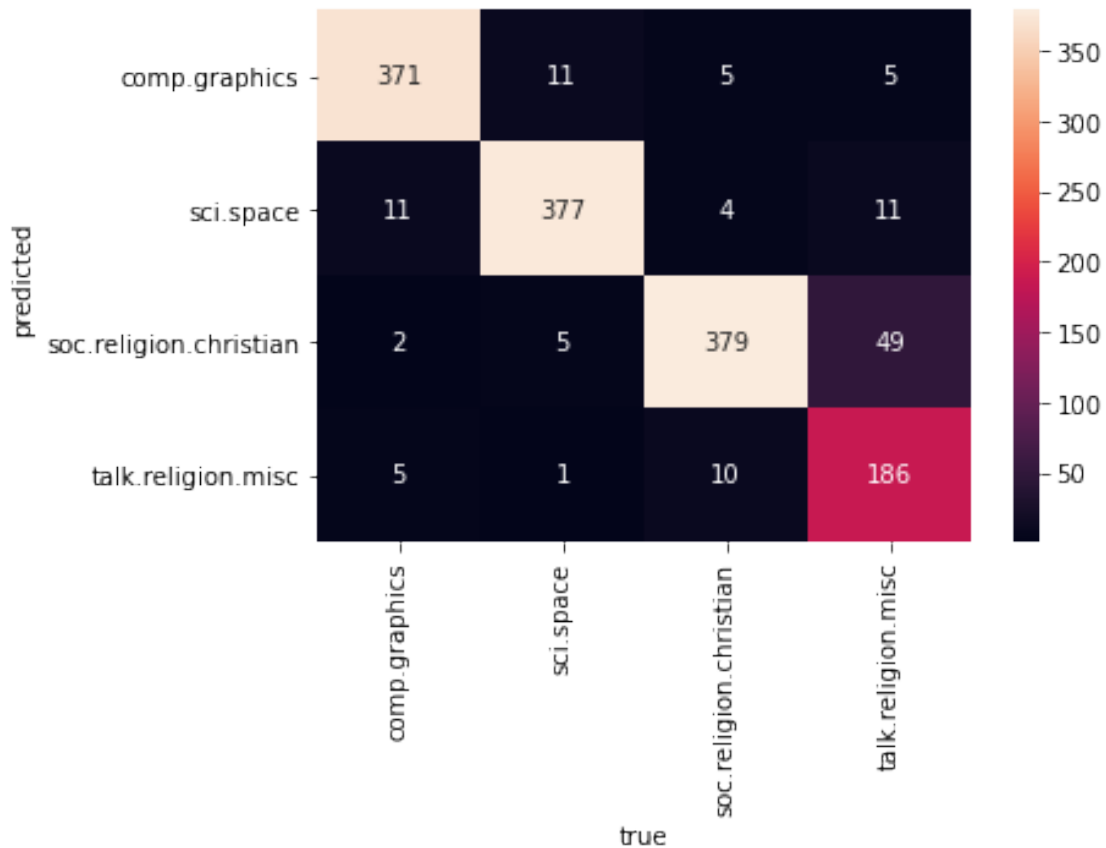
   accuracy                0.92         1432
  macro avg         0.92         0.90         0.91         1432
 weighted avg         0.92         0.92         0.92         1432
```

```
[10]: from sklearn.metrics import confusion_matrix

#use confusion matrix
mat = confusion_matrix(y_test, yhat)
```

```
import seaborn as sns
sns.heatmap(mat.T, annot=True, fmt="d",
            xticklabels=train.target_names, yticklabels=train.target_names)
plt.xlabel('true')
plt.ylabel('predicted')
```

[10]: Text(32.99999999999999, 0.5, 'predicted')



0.2 TFID

```
[11]: method = 'tfidf'
model = MyMultiNaive(method)
X_train, X_test, y_train, y_test = model.transform(train, test)
model.fit(X_train, y_train)

classes = np.unique(y_test)
yhat = model.predict(X_test, classes)
yhat = np.argmax(yhat, axis=1)

from sklearn.preprocessing import label_binarize
```

```

from sklearn.metrics import average_precision_score, classification_report

n_classes = len(np.unique(y_test))

print("Accuracy: ", np.sum(yhat == y_test)/len(y_test))

print("====Average precision score====")
y_test_binarized = label_binarize(y_test, classes=[0, 1, 2, 3])
yhat_binarized = label_binarize(yhat, classes=[0, 1, 2, 3])

for i in range(n_classes):
    class_score = average_precision_score(y_test_binarized[:, i],
    ↪yhat_binarized[:, i])
    print(f"Class {i} score: ", class_score)

print("====Classification report====")
print("Report: ", classification_report(y_test, yhat))

from sklearn.metrics import confusion_matrix

#use confusion matrix
mat = confusion_matrix(y_test, yhat)

import seaborn as sns
sns.heatmap(mat.T, annot=True, fmt="d",
            xticklabels=train.target_names, yticklabels=train.target_names)
plt.xlabel('true')
plt.ylabel('predicted')

```

Accuracy: 0.8016759776536313

====Average precision score====

Class 0 score: 0.888341920518241

Class 1 score: 0.8744630809734135

Class 2 score: 0.6122064043881043

Class 3 score: 0.332994836297269

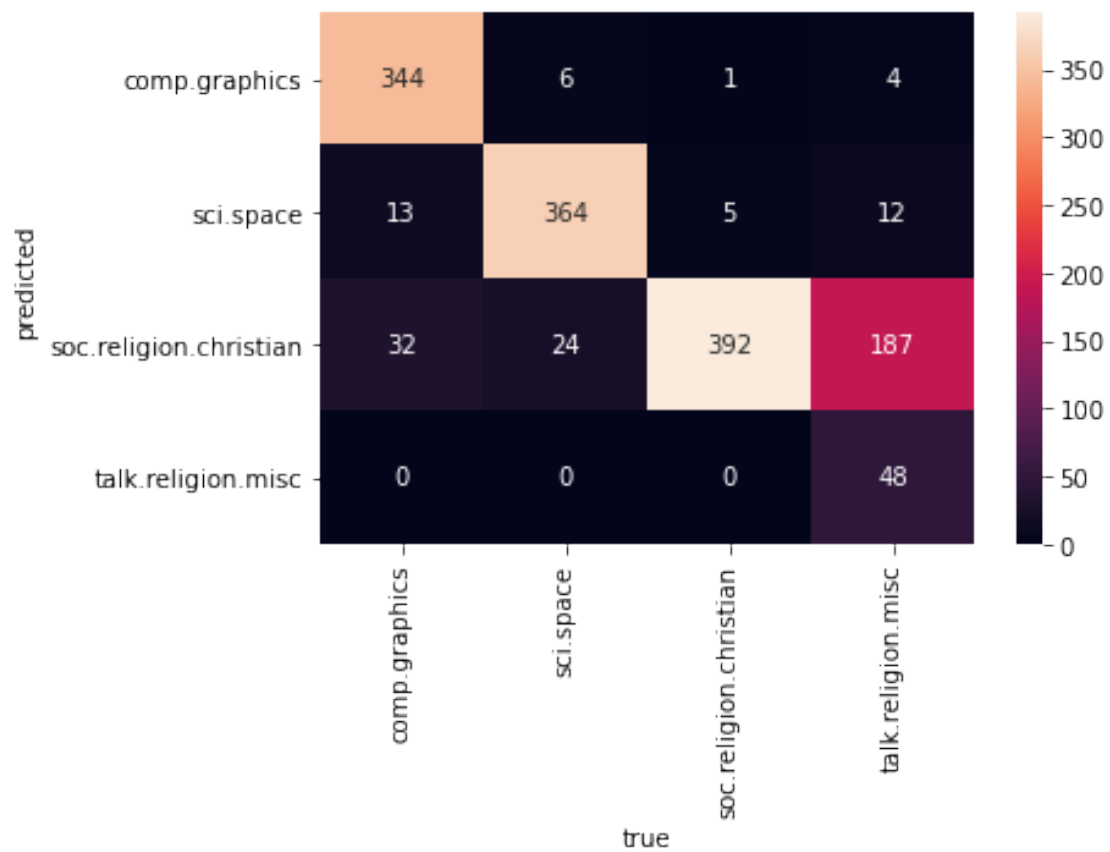
====Classification report====

Report: precision recall f1-score support

0	0.97	0.88	0.92	389
1	0.92	0.92	0.92	394
2	0.62	0.98	0.76	398
3	1.00	0.19	0.32	251

accuracy			0.80	1432
macro avg	0.88	0.75	0.73	1432
weighted avg	0.86	0.80	0.77	1432

[11]: Text(32.99999999999999, 0.5, 'predicted')



[]: