

lab_05_solution

September 16, 2021

1 Lab 5

Optimization Using Newton's Method

In this lab, we'll explore an alternative to gradient descent for nonlinear optimization problems: Newton's method.

1.1 Newton's method in one dimension

Consider the problem of finding the *roots* \mathbf{x} of a nonlinear function $f : \mathbb{R}^N \rightarrow \mathbb{R}$. A root of f is a point \mathbf{x} that satisfies $f(\mathbf{x}) = 0$.

In one dimension, Newton's method for finding zeroes works as follows:

1. Pick an initial guess x_0
2. Let $x_{i+1} = x_i + \frac{f(x_i)}{f'(x_i)}$
3. If not converged, go to #2.

Convergence occurs when $|f(x_i)| < \epsilon_1$ or when $|f(x_{i+1}) - f(x_i)| < \epsilon_2$.

Let's see how this works in practice.

```
[17]: import matplotlib.pyplot as plt
import numpy as np
from mpl_toolkits.mplot3d import Axes3D
import pandas as pd
```

```
[18]: # Example 1: Root finding for cubic polynomial

n = 200

x = np.linspace(-3, 3, n)

def fx(x, p):
    f_x = np.polyval(p, x)
    return f_x

# Create the polynomial f(x) = x^3 + x^2

p = np.poly1d([1, 1, 0, 0])
```

```

# Derivative of a polynomial
# This is a convenient method to obtain p_d = np.poly1d([3, 2, 0])

p_d = np.polyder(p)

# Get values for f(x) and f'(x) for graphing purposes

y = fx(x, p)
y_d = fx(x, p_d)

# Try three possible guesses for x0

x0_arr = [-3.0, 1.0, 3.0]
max_iter = 30
threshold = 0.001
roots = []

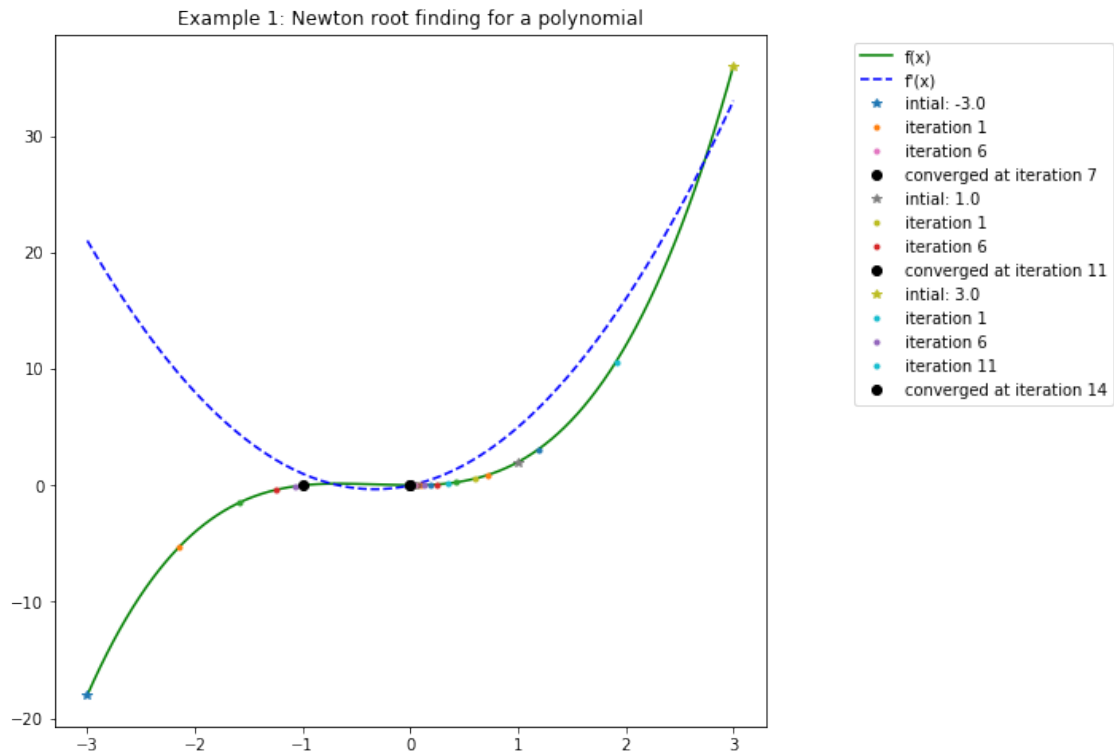
fig1 = plt.figure(figsize=(8,8))
ax = plt.axes()
plt.plot(x, y, 'g-', label='f(x)')
plt.plot(x, y_d, 'b--', label="f\'(x)")

for x0 in x0_arr:
    # Plot initial data point
    plt.plot(x0, fx(x0, p), '*', label='intial: ' + str(x0))
    i = 0
    while i < max_iter:
        #  $x_1 = x_0 - f(x_0)/f'(x_0)$ 
        x1 = x0 - fx(x0, p) / fx(x0, p_d)
        # Check for delta (x) less than threshold
        if np.abs(x0 - x1) <= threshold:
            roots.append(round(x1,4))
            break;
        # Plot current root after every 5 iterations
        if i % 5 == 0:
            plt.plot(x1, fx(x1, p), '.', label='iteration ' + str(i+1))
        else:
            plt.plot(x1, fx(x1, p), '.')
        x0 = x1
        i = i + 1
    plt.plot(x1, fx(x1, p), 'ko', label='converged at iteration ' + str(i+1))

plt.legend(bbox_to_anchor=(1.5, 1.0), loc='upper right')
plt.title('Example 1: Newton root finding for a polynomial')

plt.show()

```



[19]: *# Example 2: Root finding for sine function*

```
n = 200

x = np.linspace(-np.pi, np.pi, n)

def fx(x):
    f_x = np.sin(x)
    return f_x

def fx_d(x):
    return np.cos(x)

# Get f(x) and f'(x) for plotting

y = fx(x)
y_d = fx_d(x)

# Consider three possible starting points

x0_arr = [2.0, 1.0, -2.0]
max_iter = 30
i = 0
```

```

threshold = 0.01
roots = []

fig1 = plt.figure(figsize=(10,10))
ax = plt.axes()
ax.set_aspect(aspect = 'equal', adjustable = 'box')
plt.plot(x, y, 'g-', label='f(x)')
plt.plot(x, y_d, 'b--', label='df(x)')


for x0 in x0_arr:
    plt.plot(x0, fx(x0), 'bo', label='intial: ' + str(x0))
    i = 0;
    while i < max_iter:
        x1 = x0 - fx(x0) / fx_d(x0)
        if np.abs(x0 - x1) <= threshold:
            roots.append(x1)
            plt.plot(x1,fx(x1),'ko',label='converged at iteration '+ str(i))
            break;
        if i % 5 == 0:
            plt.plot(x1, fx(x1), '*', label='iteration '+ str(i+1))
        else:
            plt.plot(x1, fx(x1), '*')
        x0 = x1
        i = i + 1

plt.legend(bbox_to_anchor=(1.5, 1.0), loc='upper right')
plt.title('Example 2: Newton root findign for sine function')

plt.show()

print('Roots: %f, %f, %f' % (roots[0], roots[1], roots[2]))

```



Roots: 3.141593, 0.000000, -3.141593

1.2 Newton's method for optimization

Now, consider the problem of minimizing a scalar function $J : \mathbb{R}^n \mapsto \mathbb{R}$. We would like to find

$$\theta^* = \operatorname{argmin}_{\theta} J(\theta)$$

We already know gradient descent:

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \alpha \nabla_J(\theta^{(i)}).$$

But Newton's method gives us a potentially faster way to find θ^* as a zero of the system of equations

$$\nabla_J(\theta^*) = \mathbf{0}.$$

In one dimension, to find the zero of $f'(x)$, obviously, we would apply Newton's method to $f'(x)$, obtaining the iteration

$$x_{i+1} = x_i - f'(x_i)/f''(x_i).$$

The multivariate extension of Newton's optimization method is

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \mathbf{H}_f(\mathbf{x}_i) \nabla f(\mathbf{x}_i),$$

where $\mathbf{H}_f(\mathbf{x})$ is the *Hessian* of f evaluated at \mathbf{x} :

$$\mathbf{H}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 x_n} \\ \frac{\partial^2 f}{\partial x_2 x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n x_1} & \frac{\partial^2 f}{\partial x_n x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

This means, for the minimization of $J(\theta)$, we would obtain the update rule

$$\theta^{(i+1)} \leftarrow \theta^{(i)} - \mathbf{H}_J(\theta^{(i)}) \nabla_J(\theta^{(i)}).$$

1.3 Application to logistic regression

Let's create some difficult sample data as follows:

Class 1: Two features x_1 and x_2 jointly distributed as a two-dimensional spherical Gaussian with parameters

$$\mu = \begin{bmatrix} x_{1c} \\ x_{2c} \end{bmatrix}, \Sigma = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_1^2 \end{bmatrix}.$$

Class 2: Two features x_1 and x_2 in which the data are generated by first sampling an angle θ according to a uniform distribution, sampling a distance d according to a one-dimensional Gaussian with a mean of $(3\sigma_1)^2$ and a variance of $(\frac{1}{2}\sigma_1)^2$, then outputting the point

$$\mathbf{x} = \begin{bmatrix} x_{1c} + d \cos \theta \\ x_{2c} + d \sin \theta \end{bmatrix}$$

.

Generate 100 samples for each of the classes.

```
[20]: # Generate data for class 1

mu_1 = np.array([1.0, 2.0])
sigma_1 = 1
num_sample = 100
cov_mat = np.matrix([[sigma_1,0],[0,sigma_1]])
X1 = np.random.multivariate_normal(mean= mu_1, cov=cov_mat, size = num_sample)

# Generate data for class 2


angle = np.random.uniform(0, 2*np.pi, num_sample)
d = np.random.normal(np.square(3*sigma_1),np.square(.5*sigma_1),num_sample)
X2 = np.array([X1[:,0] + d*np.cos(angle), X1[:,1] + d*np.sin(angle)]).T

# Combine X1 and X2 into single dataset

X = np.concatenate([X1, X2],axis = 0)
y = np.append(np.zeros(num_sample),np.ones(num_sample))
```

```
[21]: # Plot the data

fig1 = plt.figure(figsize=(5,5))
ax = plt.axes()
plt.title('Sample data for classification problem')
plt.grid(axis='both', alpha=.25)
plt.plot(X1[:,0],X1[:,1],'b.', label = 'Class 1')
plt.plot(X2[:,0],X2[:,1],'g*', label = 'Class 2')
plt.legend(loc=2)
plt.axis('equal')
plt.show()
```



[22]: *# Split data into training and test datasets*

```
idx = np.arange(0, len(X), 1)
np.random.shuffle(idx)
idx_train = idx[0:int(.8*len(X))]
idx_test = idx[len(idx_train):len(idx)]

X_train = X[idx_train]
X_test = X[idx_test]
y_train = y[idx_train]
y_test = y[idx_test]
```

[23]: *# Normalization of data*

```
def normalization(X):
    """
    Take in numpy array of X values and return normalize X values,
    the mean and standard deviation of each feature
    """
    mean=np.mean(X,axis=0)
    std=np.std(X,axis=0)
```

```

X_norm = (X - mean)/std

return X_norm

XX = normalization(X)
X_train_norm = XX[idx_train]
X_test_norm = XX[idx_test]

```

```

[24]: # define class for logistic regression: batch gradient descent
class Logistic_BGD:

    def __init__(self):
        pass

#     def softmax(self, z):
#         z -= np.max(z)
#         return np.exp(z) / np.sum(np.exp(z))

#     def h(self, X, theta):
#         return self.softmax(np.dot(X, theta))

    def fx(self, X, theta):
        return X@theta

    def h(self, X, theta):
        # sigmoid function
        z = self.fx(X, theta)
        return 1 / (1 + np.exp(-1*z))

    def gradient(self, X, y, y_pred):
        m = len(y)
        return 1/m * np.dot(X.T, (y_pred - y))

    def costFunc(self, theta, X, y):
        m = len(y)
        y_pred = self.h(X, theta)
        error = (y * np.log(y_pred)) + ((1-y)*np.log(1-y_pred))
        cost = -1/m * sum(error)
        grad = self.gradient(X, y, y_pred)
        return cost, grad

    def gradientDescent(self, X, y, theta, alpha, num_iters):
#         X = np.insert(X, 0, 1, axis=1)
        m = len(y)
        J_history = []
        theta_history = []

```



```

    for i in range(num_iters):
        cost, grad = self.costFunc(theta,X,y)
        theta = theta - alpha * grad
        J_history.append(cost)
        theta_history.append(theta)
    J_min_index = np.argmin(J_history)
    return theta_history[J_min_index] , J_history

def predict(self,X, theta):
    labels=[]
    for i in range(X.shape[0]):
        y1=self.h(X[i], theta)
        if y1 >= 0.5:
            labels.append(1)
        else:
            labels.append(0)

    labels=np.asarray(labels)
    return labels

def checkAccuracy(self,predicted,y):
    predicted=predicted.tolist()
    y = y.tolist()
    correct=0
    for i in range(0,len(predicted)):
        if y[i]==predicted[i]:
            correct+=1
    return (float(correct)/len(predicted))*100

```

1.4 In lab exercises

1. Verify that the gradient descent solution given above is correct. Plot the optimal decision boundary you obtain.
2. Write a new class that uses Newton's method for the optimization rather than simple gradient descent.
3. Verify that you obtain a similar solution with Newton's method. Plot the optimal decision boundary you obtain.
4. Compare the number of iterations required for gradient descent vs. Newton's method. Do you observe other issues with Newton's method such as a singular or nearly singular Hessian matrix?

```

[25]: # We define a class for logistic regression (batch gradient descent)
class Logistic_NWT:

    def __init__(self):
        pass

```

```

def fx(self, X, theta):
    return X @ theta

def h(self, X, theta):
    # sigmoid function
    z = self.fx(X, theta)
    return 1 / (1 + np.exp(-1*z))

def gradient(self, X, y, y_pred):
    m = len(y)
    return 1/m * np.dot(X.T, (y_pred - y))

def hessian(self, xi, theta):
    """
    Ref : https://stats.stackexchange.com/questions/68391/hessian-of-logistic-function
    X is X_i one row of X.
    """
    y_pred = self.h(xi, theta) # sigmoid output (prediction)
    xi = xi.reshape(-1,1) # convert from 1d array to 2d (matrix)
    return xi @ xi.T * (y_pred) * (1-y_pred)

def costFunc(self, theta, X, y):
    m = len(y)
    y_pred = self.h(X, theta)
    error = (y * np.log(y_pred)) + ((1-y)*np.log(1-y_pred))
    cost = -1/m * sum(error)
    grad = self.gradient(X, y, y_pred)
    return cost, grad

def newtonMethod(self, X, y, theta, num_iters):
    m = len(y)
    J_history = []
    theta_history = []

    # see error
    found_singular_matrix = False
    error_msg = None
    for i in range(num_iters):
        hessian_mat = np.zeros((X.shape[1], X.shape[1]))

        # compute hessian matrix of x_i
        for i in range(X.shape[0]):
            hmat_xi = self.hessian(X[i], theta)

```

```

        # sum of all matrixs (by element position)
        hessian_mat += hmat_xi
    cost, grad = self.costFunc(theta, X,y)

    # replace gradient descent with newton method
    # follow the formula above.
    try :
        theta = theta - np.linalg.inv(hessian_mat) @ grad
    except Exception as e :
        error_msg = e
        found_singular_matrix = True
        # leverages SVD to approximate initial matrix.
        theta = theta - np.linalg.pinv(hessian_mat) @ grad
    J_history.append(cost)
    theta_history.append(theta)
    if found_singular_matrix :
        print('Found error while compute newton method error ! ')
        print(error_msg)
    J_min_index = np.argmin(J_history)
    return theta_history[J_min_index] , J_history

def predict(self,X, theta):
    labels=[]
    for i in range(X.shape[0]):
        y1=self.h(X[i], theta)
        if y1 >= 0.5:
            labels.append(1)
        else:
            labels.append(0)

    labels=np.asarray(labels)
    return labels

def checkAccuracy(self,predicted,y):
    predicted=predicted.tolist()
    y = y.tolist()
    correct=0
    for i in range(0,len(predicted)):
        if y[i]==predicted[i]:
            correct+=1
    return (float(correct)/len(predicted))*100

```

```

[26]: def boundary_points(X, theta):
    theta = theta.reshape(-1,1)
    v_orthogonal = np.array([[theta[1,0]], [theta[2,0]]])
    v_ortho_length = np.sqrt(v_orthogonal.T @ v_orthogonal)
    dist_ortho = theta[0,0] / v_ortho_length

```

```

v_orthogonal = v_orthogonal / v_ortho_length
v_parallel = np.array([[v_orthogonal[1,0]], [v_orthogonal[0,0]]])
projections = X @ v_parallel
proj_1 = min(projections)
proj_2 = max(projections)
point_1 = proj_1 * v_parallel - dist_ortho * v_orthogonal
point_2 = proj_2 * v_parallel - dist_ortho * v_orthogonal
return point_1, point_2

```

```

[27]: # create X to equal to polynomial function
x_df = pd.DataFrame(X, columns=['X0', 'X1'])

x_df['y'] = y

x_df['X0'] = normalization(x_df.X0)
x_df['X1'] = normalization(x_df.X1)
linX = x_df[['X0', 'X1']].values
linX = np.insert(linX, 0, 1, axis=1)

X_train = linX[idx_train]
X_test = linX[idx_test]
y_train = y[idx_train]
y_test = y[idx_test]

```

Getting optimal theta sing Logistic Batch Gradient Descent:

```

[28]: alpha = 0.05
iterations = 1000
init_theta = np.ones(X_train.shape[1])

lgx = Logistic_BGD()
theta_x, j_hist_x = lgx.gradientDescent(X_train, y_train, init_theta, alpha,
↪ iterations)

```

Getting optimal theta using Newton's Method

```

[29]: iterations = 1000
lnx = Logistic_NWT()
theta_xn, j_hist_xn = lnx.newtonMethod(X_train, y_train, init_theta, iterations)

[30]: plt.title('Cost history')
plt.plot(j_hist_xn, label='Newton')
plt.plot(j_hist_x, label='Gradient')


plt.legend()
plt.show()

```

```

print('Minimum Cost for each method')
print('Gradient :', np.min(j_hist_x))
print('Newton :', np.min(j_hist_xn))

```



Minimum Cost for each method
 Gradient : 0.6889016406221803
 Newton : 0.6889024628224568

```

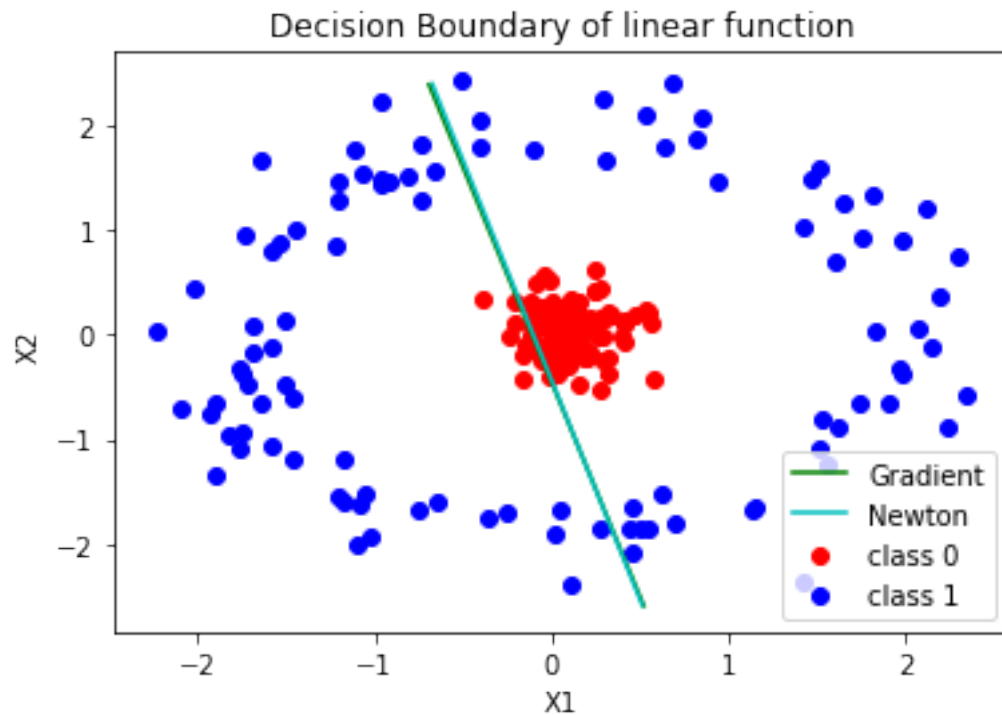
[31]: # Plotting decision boundary for the above classes from solution obtained using
      ↪ batch gradient descent and Newton's method
y0_df = x_df[x_df.y == 0]
y1_df = x_df[x_df.y == 1]

point_1, point_2 = boundary_points(linX[:,1:], theta_x)
point_1n, point_2n = boundary_points(linX[:,1:], theta_xn)

plt.title('Decision Boundary of linear function')
plt.scatter(y0_df.X0, y0_df.X1, c='r', label='class 0')
plt.scatter(y1_df.X0, y1_df.X1, c='b', label='class 1')
plt.legend()
plt.xlabel('X1')
plt.ylabel('X2')
# plot the boundaries for both methods

```

```
plt.plot([point_1[0,0], point_2[0,0]],[point_1[1,0], point_2[1,0]], 'g-',  
        ↪label='Gradient')  
plt.plot([point_1n[0,0], point_2n[0,0]],[point_1n[1,0], point_2n[1,0]], 'c-',  
        ↪label='Newton')  
plt.legend(loc=0)  
plt.show()
```



```
[32]: yg_pred = lgx.predict(X_test, theta_x)  
      yn_pred = lnx.predict(X_test, theta_xn)  
      g_acc = lgx.checkAccuracy(yg_pred, y_test)  
      n_acc = lnx.checkAccuracy(yn_pred, y_test)  
  
      print("Test Accuracy")  
      print('Gradient Accuracy : ', g_acc)  
      print('Newton Accuracy : ', n_acc)
```

```
Test Accuracy  
Gradient Accuracy : 72.5  
Newton Accuracy : 72.5
```

1.5 Take-home exercises

1. Perform a *polar transformation* on the data above to obtain a linearly separable dataset.

2. Verify that you obtain good classification accuracy for logistic regression with GD or Newton's method after the polar transformation
3. Apply Newton's method to the dataset you used for the take home exercises in Lab 03.

1.6 The report

Write a brief report covering your experiments (both in lab and take home) and send as a Jupyter notebook to the TAs, Manish and Abhishek before the next lab.

In your solution, be sure to follow instructions.

1.6.1 Performing Polar Transformations


```
[33]: # Prepare data
# Convert X to angle and radius
df = pd.DataFrame(X, columns=['X0', 'X1'])
df['angles'] = np.arctan(df.X1 / df.X0)
df['radius'] = np.sqrt(df.X0 ** 2 + df.X1 ** 2)
df['y'] = y

newX = df[['angles', 'radius']].values
newX = np.insert(newX, 0, 1, axis=1)
X_train = newX[idx_train]
X_test = newX[idx_test]
y_train = y[idx_train]
y_test = y[idx_test]
```

```
[34]: y0_df = df[df.y == 0]
y1_df = df[df.y == 1]

plt.title('Transformed Data')
plt.scatter(y0_df.angles, y0_df.radius, c='r', label='class 0')
plt.scatter(y1_df.angles, y1_df.radius, c='b', label='class 1')
plt.legend()
plt.xlabel('angle')
plt.ylabel('radius')
```

```
[34]: Text(0, 0.5, 'radius')
```



```
[35]: alpha = 0.01
iterations = 10000
init_theta = np.ones(newX.shape[1])

lg = Logistic_BGD()
theta, j_hist = lg.gradientDescent(X_train, y_train, init_theta, alpha,
    ↪ iterations)

[36]: iterations = 800
init_theta = np.ones(newX.shape[1])

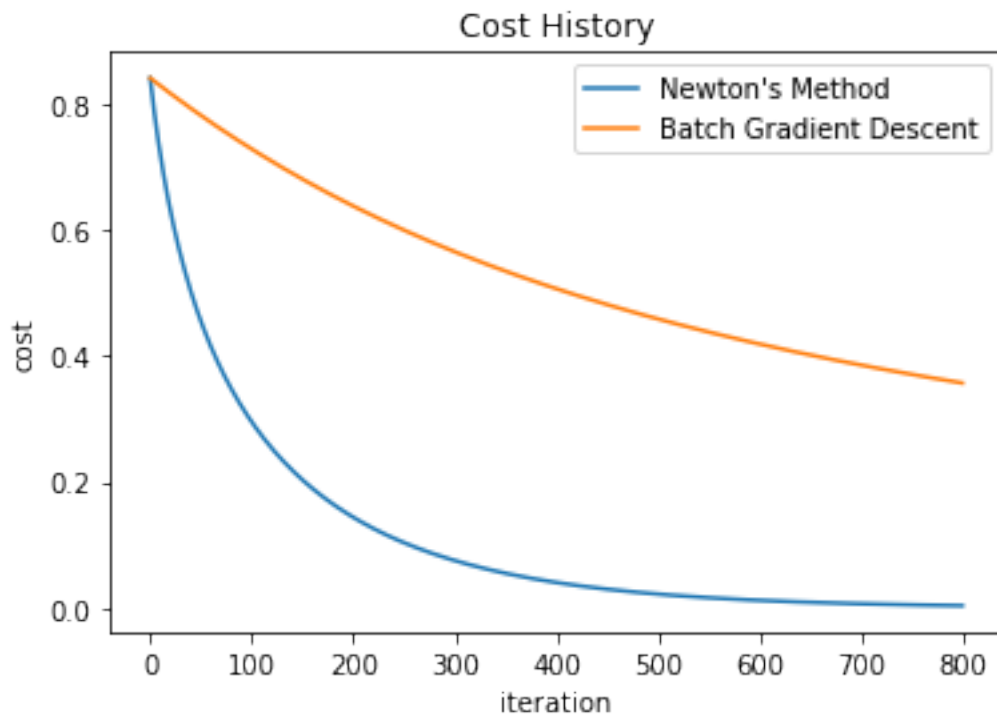
ln = Logistic_NWT()
theta_n, j_hist_n = ln.newtonMethod(X_train, y_train, init_theta, iterations)

[37]: plt.plot(j_hist_n, label='Newton\'s Method')
plt.plot(j_hist[:iterations], label='Batch Gradient Descent')
plt.title('Cost History')
plt.xlabel('iteration')
plt.ylabel('cost')
plt.legend()
plt.show()

print('Minimum Cost for each method from polar transformation')
```



```
print('Gradient Descent:', np.min(j_hist))
print('Newton's Method :', np.min(j_hist_n))
```



Minimum Cost for each method from polar transformation

Gradient Descent: 0.0449652208110658


Newton's Method : 0.0036222957878202957

```
[38]: y0_df = df[df.y == 0]
      y1_df = df[df.y == 1]

      point_1, point_2 = boundary_points(newX[:,1:], theta)
      point_1n, point_2n = boundary_points(newX[:,1:], theta_n)

      plt.title('Decision Boundary')
      plt.scatter(y0_df.angles, y0_df.radius, c='r', label='class 0')
      plt.scatter(y1_df.angles, y1_df.radius, c='b', label='class 1')
      plt.legend()
      plt.xlabel('angle')
      plt.ylabel('radius')
      # plot the boundaries for both methods
      plt.plot([point_1[0,0], point_2[0,0]], [point_1[1,0], point_2[1,0]], 'g-', lw=2,
               label='Gradient')
```

```
plt.plot([point_1n[0,0], point_2n[0,0]],[point_1n[1,0], point_2n[1,0]], 'c-',
        ↪label='Newton')
plt.legend(loc=0)
# plt.ylim(0,14)
plt.xlim(-2,2)
plt.show()
```



```
[39]: yg_pred = lg.predict(X_train, theta)
      yn_pred = ln.predict(X_train, theta_n)
      g_acc = lg.checkAccuracy(yg_pred, y_train)
      n_acc = lg.checkAccuracy(yn_pred, y_train)

      print("Train accuracy for polar transformation")
      print('Gradient Accuracy : ', g_acc)
      print('Newton Accuracy : ', n_acc)

      yg_pred = lg.predict(X_test, theta)
      yn_pred = ln.predict(X_test, theta_n)
      g_acc = lg.checkAccuracy(yg_pred, y_test)
      n_acc = lg.checkAccuracy(yn_pred, y_test)

      print("Test accuracy for polar transformation")
      print('Gradient Accuracy : ', g_acc)
```

```
print('Newton Accuracy : ', n_acc)
```

```
Train accuracy for polar transformation
Gradient Accuracy : 100.0
Newton Accuracy : 100.0
Test accuracy for polar transformation
Gradient Accuracy : 100.0
Newton Accuracy : 100.0
```

Here we can see, we get the perfect accuracy while performing polar transformation in our dataset.

1.6.2 Using Newton method in Lab 3- Loan prediction dataset

```
[40]: # Import the data

data_train = pd.read_csv('train_LoanPrediction.csv')
data_test = pd.read_csv('test_LoanPrediction.csv')

# Start to explore the data

print('Training data shape', data_train.shape)
print('Test data shape', data_test.shape)

print('Training data:\n', data_train)
```

Training data shape (614, 13)

Test data shape (367, 12)

Training data:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
..	

609	2900	0.0	71.0	360.0
610	4106	0.0	40.0	180.0
611	8072	240.0	253.0	360.0
612	7583	0.0	187.0	360.0
613	4583	0.0	133.0	360.0

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y
..
609	1.0	Rural	Y
610	1.0	Rural	Y
611	1.0	Urban	Y
612	1.0	Urban	Y
613	0.0	Semiurban	N

[614 rows x 13 columns]

```
[41]: # Now checking for the missing values of the train and test datasets

print('Missing values for train data:\n-----\n', data_train.
      ↪isnull().sum())
print('Missing values for test data \n -----\n', data_test.
      ↪isnull().sum())
```

Missing values for train data:

```
-----
Loan_ID          0
Gender           13
Married          3
Dependents       15
Education        0
Self_Employed   32
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       22
Loan_Amount_Term 14
Credit_History  50
Property_Area    0
Loan_Status      0
```

dtype: int64

Missing values for test data

```
-----
Loan_ID          0
Gender           11
```

```

Married          0
Dependents       10
Education        0
Self_Employed    23
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       5
Loan_Amount_Term 6
Credit_History  29
Property_Area    0
dtype: int64

```

```

[42]: # Filling the missing values

def fill_marital_status(data, yes_num_train, no_num_train):
    data['Married'].fillna('Yes', inplace = True, limit = yes_num_train)
    data['Married'].fillna('No', inplace = True, limit = no_num_train)

fill_marital_status(data_train, 2, 1)
fill_marital_status(data_test, 2, 1)

def fill_dependent_status(num_0_train, num_1_train, num_2_train, num_3_train,
↳ num_0_test, num_1_test, num_2_test, num_3_test):
    data_train['Dependents'].fillna('0', inplace=True, limit = num_0_train)
    data_train['Dependents'].fillna('1', inplace=True, limit = num_1_train)
    data_train['Dependents'].fillna('2', inplace=True, limit = num_2_train)
    data_train['Dependents'].fillna('3+', inplace=True, limit = num_3_train)
    data_test['Dependents'].fillna('0', inplace=True, limit = num_0_test)
    data_test['Dependents'].fillna('1', inplace=True, limit = num_1_test)
    data_test['Dependents'].fillna('2', inplace=True, limit = num_2_test)
    data_test['Dependents'].fillna('3+', inplace=True, limit = num_3_test)

fill_dependent_status(9, 2, 2, 2, 5, 2, 2, 1)

data_train['Dependents'].replace('3+', '4', inplace = True)
data_test['Dependents'].replace('3+', '4', inplace = True)

loan_amount_mean = np.mean(data_train["LoanAmount"])
data_train['LoanAmount'].fillna(loan_amount_mean, inplace=True, limit = 22)
data_test['LoanAmount'].fillna(loan_amount_mean, inplace=True, limit = 5)

# filling Gender by mode
data_train.Gender.fillna(data_train.Gender.mode()[0], inplace=True)
data_test.Gender.fillna(data_test.Gender.mode()[0], inplace=True)

# filling Self Employed by mode

```

```

data_train.Self_Employed.fillna(data_train.Self_Employed.mode()[0] ,  

    ↳inplace=True)
data_test.Self_Employed.fillna(data_test.Self_Employed.mode()[0] , inplace=True)

# filling loan amount term by mean
data_train.Loan_Amount_Term.fillna(data_train.Loan_Amount_Term.mean() ,  

    ↳inplace=True)
data_test.Loan_Amount_Term.fillna(data_test.Loan_Amount_Term.mean() ,  

    ↳inplace=True)

# filling credit history by mode
data_train.Credit_History.fillna(data_train.Credit_History.mode()[0] ,  

    ↳inplace=True)
data_test.Credit_History.fillna(data_test.Credit_History.mode()[0] ,  

    ↳inplace=True)

```

[43]: *#Now checking for the missing values of the train and test datasets*

```

print('Missing values for train data:\n-----\n', data_train.  

    ↳isnull().sum())
print('Missing values for test data \n -----\n', data_test.  

    ↳isnull().sum())

```

Missing values for train data:

```

-----
Loan_ID          0
Gender           0
Married          0
Dependents       0
Education        0
Self_Employed    0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History   0
Property_Area    0
Loan_Status      0
dtype: int64
Missing values for test data

```

```

-----
Loan_ID          0
Gender           0
Married          0
Dependents       0
Education        0
Self_Employed    0

```

```

ApplicantIncome      0
CoapplicantIncome    0
LoanAmount           0
Loan_Amount_Term     0
Credit_History       0
Property_Area        0
dtype: int64

```

```

[44]: data_test = data_test.iloc[0:,1:]
      data_train = data_train.iloc[0:,1:]

      print(data_test.info())
      print(data_train.info())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 367 entries, 0 to 366
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                367 non-null   object
1   Married               367 non-null   object
2   Dependents            367 non-null   object
3   Education             367 non-null   object
4   Self_Employed         367 non-null   object
5   ApplicantIncome       367 non-null   int64
6   CoapplicantIncome     367 non-null   int64
7   LoanAmount            367 non-null   float64
8   Loan_Amount_Term      367 non-null   float64
9   Credit_History        367 non-null   float64
10  Property_Area         367 non-null   object
dtypes: float64(3), int64(2), object(6)
memory usage: 31.7+ KB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                614 non-null   object
1   Married               614 non-null   object
2   Dependents            614 non-null   object
3   Education             614 non-null   object
4   Self_Employed         614 non-null   object
5   ApplicantIncome       614 non-null   int64
6   CoapplicantIncome     614 non-null   float64
7   LoanAmount            614 non-null   float64
8   Loan_Amount_Term      614 non-null   float64
9   Credit_History        614 non-null   float64

```

```

10 Property_Area      614 non-null    object
11 Loan_Status        614 non-null    object
dtypes: float64(4), int64(1), object(7)
memory usage: 57.7+ KB
None

```

[45]: *#Converting categorical into 0 and 1*

```

data_train['Gender'].replace('Male', 0, inplace = True)
data_train['Gender'].replace('Female', 1, inplace = True)
data_train['Married'].replace('No', 0, inplace = True)
data_train['Married'].replace('Yes', 1, inplace = True)
data_train['Education'].replace('Not Graduate', 0, inplace = True)
data_train['Education'].replace('Graduate', 1, inplace = True)
data_train['Self_Employed'].replace('No', 0, inplace = True)
data_train['Self_Employed'].replace('Yes', 1, inplace = True)
data_train['Property_Area'].replace('Urban', 0, inplace = True)
data_train['Property_Area'].replace('Semiurban', 1, inplace = True)
data_train['Property_Area'].replace('Rural', 2, inplace = True)
data_train['Loan_Status'].replace('N', 0, inplace = True)
data_train['Loan_Status'].replace('Y', 1, inplace = True)

data_test['Gender'].replace('Male', 0, inplace = True)
data_test['Gender'].replace('Female', 1, inplace = True)
data_test['Married'].replace('No', 0, inplace = True)
data_test['Married'].replace('Yes', 1, inplace = True)
data_test['Education'].replace('Not Graduate', 0, inplace = True)
data_test['Education'].replace('Graduate', 1, inplace = True)
data_test['Self_Employed'].replace('No', 0, inplace = True)
data_test['Self_Employed'].replace('Yes', 1, inplace = True)
data_test['Property_Area'].replace('Urban', 0, inplace = True)
data_test['Property_Area'].replace('Semiurban', 1, inplace = True)
data_test['Property_Area'].replace('Rural', 2, inplace = True)

print(data_train.info())
print(data_test.info())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                614 non-null   int64
1   Married               614 non-null   int64
2   Dependents            614 non-null   object
3   Education              614 non-null   int64
4   Self_Employed         614 non-null   int64

```



```

5  ApplicantIncome      614 non-null    int64
6  CoapplicantIncome    614 non-null    float64
7  LoanAmount           614 non-null    float64
8  Loan_Amount_Term     614 non-null    float64
9  Credit_History       614 non-null    float64
10 Property_Area        614 non-null    int64
11 Loan_Status          614 non-null    int64

```

dtypes: float64(4), int64(7), object(1)

memory usage: 57.7+ KB

None

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 367 entries, 0 to 366

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Gender	367 non-null	int64
1	Married	367 non-null	int64
2	Dependents	367 non-null	object
3	Education	367 non-null	int64
4	Self_Employed	367 non-null	int64
5	ApplicantIncome	367 non-null	int64
6	CoapplicantIncome	367 non-null	int64
7	LoanAmount	367 non-null	float64
8	Loan_Amount_Term	367 non-null	float64
9	Credit_History	367 non-null	float64
10	Property_Area	367 non-null	int64

dtypes: float64(3), int64(7), object(1)

memory usage: 31.7+ KB

None

```

[46]: print(data_train.dtypes)
data_train['Dependents'] = data_train['Dependents'].astype(str).astype(int)
data_train['Loan_Amount_Term'] = data_train['Loan_Amount_Term'].astype(int)
data_train['Credit_History'] = data_train['Credit_History'].astype(int)
data_train['CoapplicantIncome'] = data_train['CoapplicantIncome'].astype(int)
data_train['LoanAmount'] = data_train['LoanAmount'].astype(int)
print(data_train.dtypes)

```

```

Gender                int64
Married               int64
Dependents            object
Education             int64
Self_Employed        int64
ApplicantIncome       int64
CoapplicantIncome     float64
LoanAmount            float64
Loan_Amount_Term     float64
Credit_History       float64

```

```

Property_Area      int64
Loan_Status        int64
dtype: object
Gender             int64
Married            int64
Dependents         int32
Education          int64
Self_Employed      int64
ApplicantIncome    int64
CoapplicantIncome  int32
LoanAmount         int32
Loan_Amount_Term   int32
Credit_History     int32
Property_Area      int64
Loan_Status        int64
dtype: object

```

```
[47]: data_train
```

```

[47]:      Gender  Married  Dependents  Education  Self_Employed  ApplicantIncome  \
0         0         0           0           1           0           5849
1         0         1           1           1           0           4583
2         0         1           0           1           1           3000
3         0         1           0           0           0           2583
4         0         0           0           1           0           6000
..      ...      ...      ...      ...      ...      ...
609        1         0           0           1           0           2900
610        0         1           4           1           0           4106
611        0         1           1           1           0           8072
612        0         1           2           1           0           7583
613        1         0           0           1           1           4583

      CoapplicantIncome  LoanAmount  Loan_Amount_Term  Credit_History  \
0                   0         146           360           1
1                1508         128           360           1
2                   0          66           360           1
3                2358         120           360           1
4                   0         141           360           1
..      ...      ...      ...      ...
609                   0          71           360           1
610                   0          40           180           1
611                240         253           360           1
612                   0         187           360           1
613                   0         133           360           0

      Property_Area  Loan_Status
0                 0             1

```

1	2	0
2	0	1
3	0	1
4	0	1
..
609	2	1
610	2	1
611	0	1
612	0	1
613	1	0

[614 rows x 12 columns]

```
[48]: # apply newton method with Lab 3 dataset
iterations = 1000
init_theta = np.ones(X_train.shape[1])

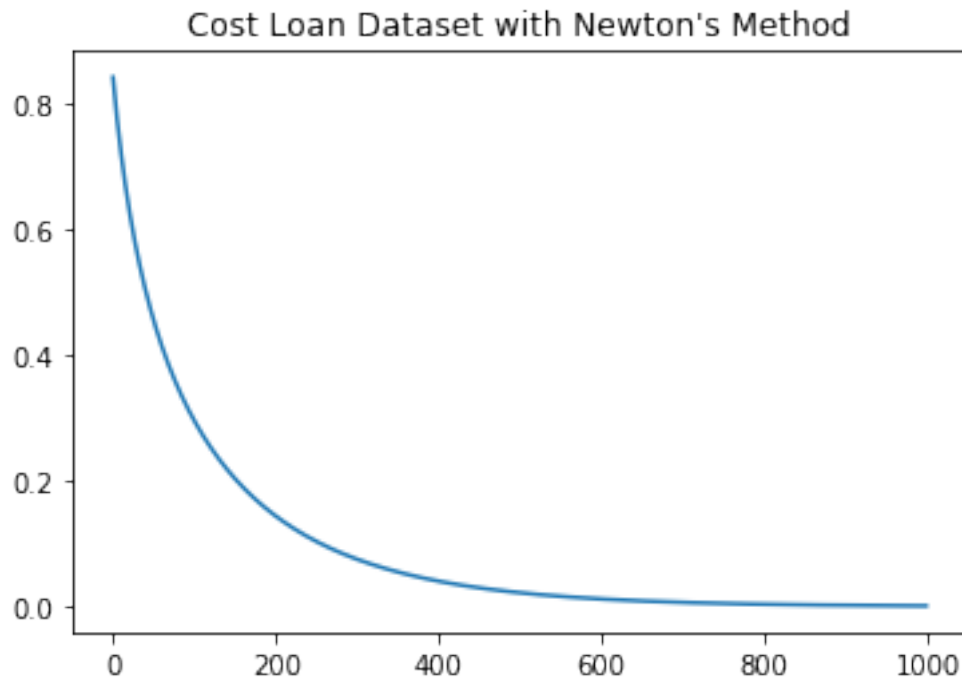
ln_lab = Logistic_NWT()
theta_nlab, j_hist_nlab = ln_lab.newtonMethod(X_train, y_train, init_theta,
→iterations)
```

```
[49]: plt.title('Cost Loan Dataset with Newton\'s Method')
plt.plot(j_hist_nlab)
plt.show()

yn_pred = ln_lab.predict(X_train, theta_nlab)
n_acc = ln_lab.checkAccuracy(yn_pred, y_train)
print("Train Set Accuracy From Lab3 dataset")
print('Newton Accuracy : ', n_acc)

yn_pred = ln_lab.predict(X_test, theta_nlab)
n_acc = ln_lab.checkAccuracy(yn_pred, y_test)

print("Test Set Accuracy From Lab3 dataset")
print('Newton Accuracy : ', n_acc)
```



Train Set Accuracy From Lab3 dataset
Newton Accuracy : 100.0
Test Set Accuracy From Lab3 dataset
Newton Accuracy : 100.0

[]:

[]: